

Newspaper Article Classification Contest

Final Report

Terry Li

tianweil@andrew.cmu.edu

Tao Yu

taoyu@andrew.cmu.edu

Abstract

This is the final report for the Newspaper Article Classification Contest project. The content includes how we retrieve the features from articles, which classifier we used, how to train the classifier and the performance of our classifier.

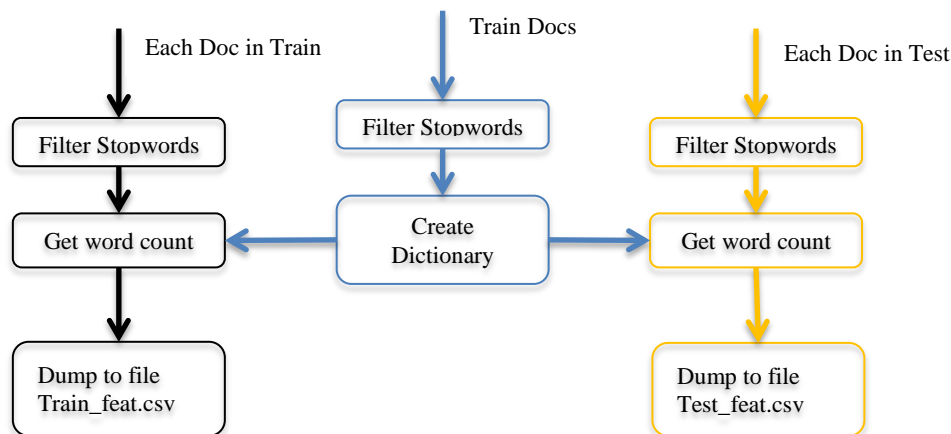
1 Project members

Terry Li (Andrew ID: tianweil) and Tao Yu (Andrew ID: taoyu)

2 Feature Extraction

2.1 Basic Approach

We use the word count as the features of our classifier. The process to extract features for train articles and test articles is as below.



For each article:

1. Parse the articles and separate it to a bunch of words
2. Remove the stop words
3. Count the word count for each word in the dictionary for each article
4. Use the word count as the features for classification

After filtering the stop words, we have 38800 words left. Then we use these words to classify the articles.

2.2 Word Stemming

To reduce the number of features, we leveraged the public Porter Stemming Algorithm [6] to pre-process the documents. After stemming, the number features reduced to 25313.

2.3 TF-IDF

To reduce the number of features, we also implemented the tf-idf feature selection algorithm. The original tf-idf algorithm does not take the class information into account. It only calculates the score of a certain word in a bunch of documents. The original equation is

$$idf(word) = \log\left(\frac{doc\ count}{doc\ count\ that\ contains\ word}\right)$$

To introduce the class information into this equation, we used a modified version of tf-idf:

$$idf(word|c) = \log\left(\frac{doc\ count * doc\ count\ in\ class\ c\ and\ contains\ word}{doc\ count\ that\ contains\ word}\right)$$

For the tf, the equation is the same:

$$tf(word|c) = \frac{word\ count\ in\ class\ c}{all\ word\ count\ in\ class\ c}$$

Because the word count for a certain word in a certain class could be 0, and the doc count that contains a certain word in certain class also could be 0, we introduced add-1 smooth in the process of calculating both tf and idf.

After calculating tf-idf for every feature, we selected 2500 features with the highest scores to be the features of our classifier.

3 Naïve Bayes Classifier

For Naïve Bayes classifier,

$$C_{NB} = \operatorname{argmax} P(x_1, x_2, \dots, x_n|c)P(c) = \operatorname{argmax} P(c) \prod P(x_i|c)$$

We use the word count of each to calculate the $P(x_i|c)$ for every class c .

In the beginning, the equation we used is:

$$P(x_i|c) = \frac{N_{ci}}{N_c}$$

In which N_{ci} is the number of times word x_i appears in the documents in class c . N_c is the total number of words appears in documents in class c .

Because a certain word may have zero count in a class, then the $P(x_i|c)$ is 0. This will cause the posterior probability will become 0. Because there are always some words have 0 count in some classes. It makes all posterior probability for every class is 0, which makes our classifier useless.

To solve this problem, we introduced add-1 smooth to the probability calculation. That is we add 1 count for every feature (different word). To make the total probability of all feature still 1, we need to add the number of features to the total word count.

70 Then the equation becomes:

71

$$P(x_i | c) = \frac{N_{ci} + 1}{N_c + n}$$

72

73 In which n is the number of features.

74

75 For the Prior probability of each class, we just use MLE to calculate them.

76

77 After get the $P(c)$ and $P(x_i|c)$ for every word and every class, we finished training the
78 classifier.

79

80

81 **4 Multinomial Naive Bayes**

82

83 Multinomial Naive Bayes models the distribution of words in a document as a multinomial.
84 The likelihood of a document is a product of the probability of the words that appear in the
85 document.

86

87 So the probability of one document belongs to class c is:

88

$$P(c|doc) = P(c) \prod P(x_i|c)^{f_i}$$

89

90 Where the f_i is the frequency of word i in doc.

91

92 To simplify the calculation, we use the log likelihood instead of original probability.

93

$$L(c|doc) = \log P(c) + \sum_{i=0}^n f_i * \log P(x_i|c)$$

94

95 Then we can classify the doc into the class that has the highest $L(c|doc)$

96

97

98

99 **5 Modified Naive Bayes Classifier**

100

101 According to paper [5], the performance of Naive Bayes Classifier is limited by the prior
102 probability $P(c)$ and the probability of word i in each class c $P(x_i|c)$. We tried the
103 modification introduced to Naive Bayes Classifier introduced by paper [5].

104

105 **5.1 Complement Naive Bayes (CNB)**

106

107 Instead of calculating the probability of word i in class c:

$$P(x_i | c) = \frac{N_{ci} + 1}{N_c + n}$$

108

109 We calculate the probability of word i that is not in class c:

$$P(x_i | \bar{c}) = \frac{N_{\bar{c}i} + 1}{N_{\bar{c}} + n}$$

110

111 Then the log likelihood of doc belongs to class c is:

$$L(c|doc) = \log P(c) - \sum_{i=0}^n f_i * \log P(x_i|\tilde{c})$$

5.2 Transformed Naïve Bayes (TNB)

Instead of using of $P(x_i|c)$ as the probability of word i belong to class c , we can also use the tf-idf score of the word i .

Then the log likelihood of doc belongs to class c is:

$$L(c|doc) = \log P(c) + \sum_{i=0}^n f_i * \text{score}_{\text{tfidf}}(i, c)$$

6 Support Vector Machine

We implement a simple version of support vector machine algorithm with kernel parameters of linear, Poly and RBF. For convenience, we choose Quadprog as dual problem solver. We do experiments on linear SVM as well as with RBF kernel as below:

6.1 Linear SVM

SVM are shown to handle feature redundancy well, because of the reason that we have 38863 features. It is reasonable to use linear SVM. For features, we use wordcount and TF-IDF features to train the linear SVM model. Every feature will contribute to the improvement of the linear SVM model.

C is essentially a regularization parameter, which controls the trade-off between achieving a low error on the training data and minimizing the norm of the weights.

The C parameter tells the SVM optimization how much we want to avoid misclassifying each training example. For large values of C , the optimization will choose a smaller-margin hyperplane if that hyperplane does a better job of getting all the training points classified correctly. Conversely, a very small value of C will cause the optimizer to look for a larger-margin separating hyperplane, even if that hyperplane misclassifies more points. For very tiny values of C , we should get misclassified examples, often even if our training data is linearly separable.

We adjust C from $2^{(-10)}$ to 2^5 , and find that the accuracy is highest when C is $2^{(-10)}$.

6.2 SVM with RBF kernel

For SVM with RBF kernel, we replace natural product with kernel function. So, at the end, w^* would look like,

$$w^* = \sum_{i \in SV} h_i y_i \Phi(x_i)$$

and hence,

$$\langle w^*, \Phi(x) \rangle = \sum_{i \in SV} h_i y_i \langle \Phi(x_i), \Phi(x) \rangle$$

158 Similarly,

$$b^* = \frac{1}{|SV|} \sum_{i \in SV} \left(y_i - \sum_{j=1}^N (h_j y_j \langle \Phi(x_j), \Phi(x_i) \rangle) \right)$$

159

160

161 and our classification looks like

$$c_x = \text{sign}(\langle w, \Phi(x) \rangle + b)$$

162

163

164

165 After fine-tuned our learner parameters on evaluation data, we test our fitted SVM's
166 performance on the testing data that's previously unseen in the training and fine-tuning stages.
167 Because of time limitation, we use the grid search: Set a range of feasible values for C, for
168 instance C in [2, 2^2, 2^3, 2^4, 2^5, 2^6]. Set a range of feasible values for Gamma [1, 10,
169 100, 1000]. Then we use grid search to find the best C and Gamma.

170

171

172

173 6.3 SVM Experimental Result

174

Feature Vector Kernel	Kernel Type	Prediction Success Rate	Kernel Variable
wordcount	Linear	66.88%	C = 2^(-10)
TFIDF	Linear	66.94%	C = 2^(-10)
wordcount	RBF	67.01%	G = 0.5
TFIDF	RBF	67.05%	G = 0.5

175

176 Compare wordcount with TFIDF features, TFIDF is a little better than original wordcount
177 information.

178

179 Alternatively, we can use n-cross validation to estimate our SVM's performance. If we have a
180 limited amount of annotated texts, n-cross validation is recommended as it takes advantage
181 of using all the data we have.

182

183 The experimental results show that SVM consistently achieve performance on text
184 classification tasks. SVMs eliminate the need for feature selection, making the application of
185 text classification easier. Meanwhile, SVM has the advantage of robustness.

186

187

188

189

190

191

192

193

194

195

196

197

198

199


200

201

202

7 Performance

Feature selection method	Accuracy
Naïve Bayes with stemming	64.55%
Transformed Naïve Bayes	65.3%
Complement Naïve Bayes	66.7%
Basic Naïve Bayes	66.75%
SVM	67.05%



Performance Analysis:

Based on our classification result, Naïve Bayes Classifier does not have overfitting problem, so the more features the better. That's why the basic Multinomial Naïve Bayes Classifier has a better performance than Naïve Bayes with tf-idf.

Dealing with text classification, we have to consider very many (more than 30000) features. SVMs use overfitting protection, which has the potential to handle these large feature spaces.

Although there are original 38663 features, chances are that features ranked lowest still contain useful information. As a result, aggressive feature selection may lead to information loss. Naïve Bayes and SVM with all the features achieve good performance.

When we convert texts into vector representations, they contain only few entries which are not zero. Thus the vectors are very sparse, SVMs is well suited for problems with sparse instances. Meanwhile, most of the text classification problems are linearly separable. SVMs with (linear, RBF) is good at finding such a separator.

228 **References**

- 229 [1] http://en.wikipedia.org/wiki/Latent_Dirichlet_allocation
230 [2] <http://nlp.stanford.edu/IR-book/html/htmledition/feature-selection-1.html>
231 [3] Joachims, Thorsten. Text categorization with support vector machines: Learning with
232 many relevant features. Springer Berlin Heidelberg, 1998.
233 [4] Yang, Yiming, and Jan O. Pedersen. "A comparative study on feature selection in text
234 categorization." ICML. Vol. 97. 1997.
235 [5] Jason D. M. Rennie, Lawrence Shih, Jaime Teevan, David R. Karger "Tackling the Poor
236 Assumptions of Naive Bayes Text Classifiers"
237 [6] <http://tartarus.org/martin/PorterStemmer/index-old.html>
238 [7] http://www.cs.cornell.edu/people/tj/publications/joachims_98a.pdf
239 [8] <http://www.cs.cornell.edu/People/tj/svmcatbook/>